

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [43]: df=pd.read_csv('bank_nifty.csv')
```

```
In [44]: df.head()
```

```
Out[44]:
```

	Date	Open	High	Low	Close	Volume
0	1/1/2018	25565.75	25588.00	25271.55	25318.10	57576913
1	1/2/2018	25382.20	25425.50	25232.80	25338.25	72033811
2	1/3/2018	25425.75	25454.90	25300.90	25318.60	59730356
3	1/4/2018	25367.65	25490.35	25310.30	25462.60	105995860
4	1/5/2018	25524.45	25643.35	25499.55	25601.85	123622612

```
In [45]: df.describe()
```

```
Out[45]:
```

	Open	High	Low	Close
count	955.000000	955.000000	955.000000	955.000000
mean	28575.225969	28805.647958	28291.096178	28550.485812
std	4823.393891	4813.805396	4831.728274	4827.044525
min	16759.950000	17681.700000	16116.250000	16917.650000
25%	25487.200000	25651.575000	25252.175000	25443.475000
50%	27972.950000	28185.150000	27777.400000	28021.700000
75%	31523.500000	31742.175000	31240.750000	31515.000000
max	41234.550000	41829.600000	40829.150000	41238.300000

In [46]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 955 entries, 0 to 954
Data columns (total 6 columns):
 #   Column  Non-Null Count  Dtype  
---  -
 0   Date    955 non-null    object  
 1   Open    955 non-null    float64  
 2   High    955 non-null    float64  
 3   Low     955 non-null    float64  
 4   Close   955 non-null    float64  
 5   Volume  955 non-null    object  
dtypes: float64(4), object(2)
memory usage: 44.9+ KB
```

In [47]: `df.isna().sum()` *#no missing values*

```
Out[47]: Date      0
Open      0
High      0
Low       0
Close     0
Volume    0
dtype: int64
```

In [80]: `df.corr()`

```
Out[80]:
```

	Open	High	Low	Close
Open	1.000000	0.998461	0.998387	0.996952
High	0.998461	1.000000	0.997471	0.998621
Low	0.998387	0.997471	1.000000	0.998409
Close	0.996952	0.998621	0.998409	1.000000

All the variables have a very high correlation with each other

```
In [81]: x=df['High']  
y=df['Close']
```

Linear Regression Model

```
In [82]: from sklearn.model_selection import train_test_split  
x_train, x_test, y_train, y_test = train_test_split(x, y)
```

```
In [83]: from sklearn.linear_model import LinearRegression
```

```
In [84]: model = LinearRegression()
```

```
In [85]: m=model.fit(np.array(x_train).reshape(len(np.array(x_train)), -1),y_train)
```

```
In [90]: y_pred=m.predict(np.array(x_test).reshape(len(np.array(x_test)), -1))
```

```
In [91]: from sklearn.metrics import r2_score  
r2_score(y_test,y_pred)
```

```
Out[91]: 0.9964581228659305
```

The model is a very good fit giving an accuracy of 99.6458%

Multiple Linear Model

```
In [129]: x=df.iloc[:,1:4]  
y=df['Close']
```

```
In [130]: from sklearn.model_selection import train_test_split  
x_train, x_test, y_train, y_test = train_test_split(x, y)
```

```
In [132]: model = LinearRegression().fit(np.array(x_train).reshape(len(np.array(x_train)), -1), y_train)
```

```
In [134]: y_pred=model.predict(np.array(x_test).reshape(len(np.array(x_test)), -1))
```

```
In [135]: from sklearn.metrics import r2_score  
r2_score(y_test, y_pred)
```

```
Out[135]: 0.9991564063177207
```

The model has an accuracy of 99.9156% and is a better fit than linear regression model

Taking a new dataset

```
In [75]: data=pd.read_csv('adult.csv')
```

In [76]: `data.head()`

Out[76]:

	age	workclass	fnlwgt	education	educational-num	marital-status	occupation	relationship	race	gender	capital-gain	capital-loss	hours-per-week	native-country
0	25	Private	226802	11th	7	Never-married	Machine-op-inspct	Own-child	Black	Male	0	0	40	United-States
1	38	Private	89814	HS-grad	9	Married-civ-spouse	Farming-fishing	Husband	White	Male	0	0	50	United-States
2	28	Local-gov	336951	Assoc-acdm	12	Married-civ-spouse	Protective-serv	Husband	White	Male	0	0	40	United-States
3	44	Private	160323	Some-college	10	Married-civ-spouse	Machine-op-inspct	Husband	Black	Male	7688	0	40	United-States
4	18	?	103497	Some-college	10	Never-married	?	Own-child	White	Female	0	0	30	United-States

In [77]: `data.isnull().sum()` *#No null values*

Out[77]:

```

age                0
workclass          0
fnlwgt             0
education          0
educational-num    0
marital-status     0
occupation         0
relationship       0
race              0
gender            0
capital-gain       0
capital-loss       0
hours-per-week     0
native-country     0
income            0
dtype: int64

```

```
In [78]: data.info()
```

```
data columns (total 15 columns):
#   Column                Non-Null Count  Dtype
---  -
0   age                   48842 non-null  int64
1   workclass              48842 non-null  object
2   fnlwgt                 48842 non-null  int64
3   education              48842 non-null  object
4   educational-num        48842 non-null  int64
5   marital-status         48842 non-null  object
6   occupation             48842 non-null  object
7   relationship           48842 non-null  object
8   race                   48842 non-null  object
9   gender                 48842 non-null  object
10  capital-gain            48842 non-null  int64
11  capital-loss            48842 non-null  int64
12  hours-per-week          48842 non-null  int64
13  native-country         48842 non-null  object
14  income                 48842 non-null  object
dtypes: int64(6), object(9)
memory usage: 5.6+ MB
```

```
In [79]: data['income']=data['income'].map({data.income.unique()[0]:0,data.income.unique()[1]:1})
```


```
In [82]: data['gender']=data['gender'].map({data.gender.unique()[0]:1,data.gender.unique()[1]:0})
```

```
In [83]: data.gender.unique()[0]
```

```
Out[83]: 1
```

```
In [84]: sns.heatmap(data.corr())
```

```
Out[84]: <AxesSubplot:>
```



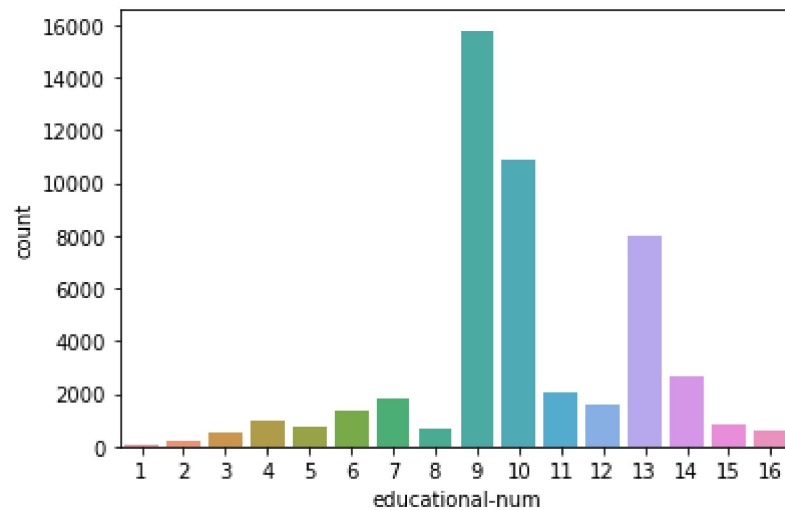
We can observe from the data that income is weekly correlated with gender. Hence the variable doesn't have a major impact on the variable.

```
In [6]: sns.countplot(data['educational-num'])
```

C:\Users\Admin\Anaconda3\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

```
warnings.warn(
```

```
Out[6]: <AxesSubplot:xlabel='educational-num', ylabel='count'>
```




```
In [7]: d=pd.get_dummies(data,drop_first=True)
d.head()
```

Out[7]:

native- y_South	native- country_Taiwan	native- country_Thailand	native- country_Trinidad&Tobago	native- country_United- States	native- country_Vietnam	native- country_Yugoslavia	income_>50K
0	0	0	0	1	0	0	0
0	0	0	0	1	0	0	0
0	0	0	0	1	0	0	1
0	0	0	0	1	0	0	1
0	0	0	0	1	0	0	0



```
In [52]: d.isnull().sum()  #no missing values
```

```
Out[52]: age                0
fnlwgt                    0
educational-num          0
capital-gain              0
capital-loss              0
..
native-country_Trinidad&Tobago  0
native-country_United-States    0
native-country_Vietnam          0
native-country_Yugoslavia       0
income_>50K                     0
Length: 101, dtype: int64
```

Logistic Model

```
In [14]: from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report, confusion_matrix
model = LogisticRegression(random_state=0)
```

```
In [32]: x=d.iloc[:,0:100]
y=d.iloc[:, -1]
```

```
In [36]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.33,random_state=42)
```

```
In [38]: m2=model.fit(x_train, y_train)
```

```
In [40]: y_pr=m2.predict(x_test)
```

```
In [43]: confusion_matrix(y_test,y_pr)
```

```
Out[43]: array([[11942,   374],
               [ 2785,  1017]], dtype=int64)
```

```
In [48]: classification_report(y_test,y_pr)
```

```
Out[48]: '
           precision    recall  f1-score   support\n\n
0           0.81      0.97      0.88       1231
6\n          1         0.73      0.27      0.39       3802\n\n
118\n  macro avg       0.77      0.62      0.64       16118\n\n
118\n  weighted avg       0.79      0.80      0.77       16118\n'
```

```
In [51]: (11942+1017)/((11942+374+2785+1017))
```

```
Out[51]: 0.8040079414319394
```

Hence the accuracy of the logistic regression model is 80.4%

